

▼ mnist

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
%matplotlib notebook
import matplotlib.pyplot as plt
from keras.initializers import he_normal
```

↳ Using TensorFlow backend.

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %t

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

↳ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```
# to get number of train and test points
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d, %d)" % X_train.shape[1:], X_train.shape[1:])
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d, %d)" % X_test.shape[1:], X_test.shape[1:])
```

↳ Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
# after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"):
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"):
```

```
↳ Number of training examples : 60000 and each image is of shape (784)
   Number of training examples : 10000 and each image is of shape (784)
```

```
# An example data point
print(X_train[0])
```

```
↳ [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
82 56 39  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190 2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 249 253 249 64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
```

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train=X_train[:255]
```

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

```
↳ Class label of first image : 5
   After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

```
# some model parameters

output_dim = 10 # output layer
input_dim = X_train.shape[1]# 784

batch_size = 128 # batch size
nb_epoch = 20# epoch
```

```
X_train.shape[1]
```

```
↳ 784
```

▼ 2 layer without dropout and batchnormal

```
#he_normal(seed=None)
model_relu = Sequential()
model_relu.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
model_relu.add(Dense(224, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_relu.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
```

```
↳
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 131us/step - loss: 0.2231 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0821 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0523 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0345 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0253 - acc: 0.

Epoch 6/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0203 - acc: 0.

Epoch 7/20

60000/60000 [=====] - 7s 118us/step - loss: 0.0156 - acc: 0.

Epoch 8/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0173 - acc: 0.

Epoch 9/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0141 - acc: 0.

Epoch 10/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0121 - acc: 0.

Epoch 11/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0122 - acc: 0.

Epoch 12/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0100 - acc: 0.

Epoch 13/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0092 - acc: 0.

Epoch 14/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0109 - acc: 0.

Epoch 15/20

60000/60000 [=====] - 7s 119us/step - loss: 0.0083 - acc: 0.

Epoch 16/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0061 - acc: 0.

Epoch 17/20

60000/60000 [=====] - 7s 120us/step - loss: 0.0105 - acc: 0.

Epoch 18/20

60000/60000 [=====] - 7s 122us/step - loss: 0.0057 - acc: 0.

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# print(history.history.keys())
```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
```

```
# we will get val_loss and val_acc only when you pass the paramter validation_data
```

```
# val_loss : validation loss
```


```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
# for each key in history.history we will have a list of length equal to number of epoch
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07829980347480105
Test accuracy: 0.9832



2 layer without drop out with batch norm

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(500, activation='relu', input_shape=(input_dim,)), kernel_initializer=
model_batch.add(BatchNormalization())

model_batch.add(Dense(224, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbos
```

↳

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 153us/step - loss: 0.1852 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0686 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 8s 138us/step - loss: 0.0445 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0340 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 8s 139us/step - loss: 0.0262 - acc: 0.

Epoch 6/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0214 - acc: 0.

Epoch 7/20

60000/60000 [=====] - 8s 137us/step - loss: 0.0195 - acc: 0.

Epoch 8/20

60000/60000 [=====] - 8s 135us/step - loss: 0.0179 - acc: 0.

Epoch 9/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0158 - acc: 0.

Epoch 10/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0134 - acc: 0.

Epoch 11/20

60000/60000 [=====] - 8s 137us/step - loss: 0.0136 - acc: 0.

Epoch 12/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0097 - acc: 0.

Epoch 13/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0103 - acc: 0.

Epoch 14/20

60000/60000 [=====] - 8s 134us/step - loss: 0.0110 - acc: 0.

Epoch 15/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0112 - acc: 0.

Epoch 16/20

60000/60000 [=====] - 8s 136us/step - loss: 0.0070 - acc: 0.

Epoch 17/20

60000/60000 [=====] - 8s 135us/step - loss: 0.0090 - acc: 0.

Epoch 18/20

60000/60000 [=====] - 8s 133us/step - loss: 0.0057 - acc: 0.

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# print(history.history.keys())
```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
```

```
# we will get val_loss and val_acc only when you pass the paramter validation_data
```

```
# val_loss : validation loss
```


```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
# for each key in history.history we will have a list of length equal to number of epoch
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0779731295770489
Test accuracy: 0.982



2 layer with drop out without batch norm

```
model_drop = Sequential()

model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,)), kernel_initializer=
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(224, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
```

↳

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 9s 149us/step - loss: 0.4279 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1877 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 8s 132us/step - loss: 0.1456 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1229 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1114 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0977 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0875 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 8s 131us/step - loss: 0.0824 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0759 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0690 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0679 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0668 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0634 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0593 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0546 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0557 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0528 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0526 - acc: 0.
```

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
```



```
# for each key in history.history we will have a list of length equal to number of epoch
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06712127507654286
Test accuracy: 0.9829



▼ 2 layer with dropout and batch norm

```
model_drop = Sequential()

model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,)), kernel_initializer=
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(224, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 175us/step - loss: 0.4112 - acc: 0

Epoch 2/20

60000/60000 [=====] - 9s 150us/step - loss: 0.1951 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 9s 150us/step - loss: 0.1535 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 9s 149us/step - loss: 0.1283 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 9s 150us/step - loss: 0.1185 - acc: 0.

Epoch 6/20

60000/60000 [=====] - 9s 150us/step - loss: 0.1025 - acc: 0.

Epoch 7/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0984 - acc: 0.

Epoch 8/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0896 - acc: 0.

Epoch 9/20

60000/60000 [=====] - 9s 149us/step - loss: 0.0837 - acc: 0.

Epoch 10/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0815 - acc: 0.

Epoch 11/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0732 - acc: 0.

Epoch 12/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0698 - acc: 0.

Epoch 13/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0679 - acc: 0.

Epoch 14/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0643 - acc: 0.

Epoch 15/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0645 - acc: 0.

Epoch 16/20

60000/60000 [=====] - 9s 150us/step - loss: 0.0572 - acc: 0.

Epoch 17/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0584 - acc: 0.

Epoch 18/20

60000/60000 [=====] - 9s 153us/step - loss: 0.0568 - acc: 0.

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# print(history.history.keys())
```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
```

```
# we will get val_loss and val_acc only when you pass the paramter validation_data
```

```
# val_loss : validation loss
```


```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
# for each key in history.history we will have a list of length equal to number of epoch
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.057800181532799616
Test accuracy: 0.9826



▼ three layered mlp relu activation

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(500, activation='relu', input_shape=(input_dim,)))# first layer 500
model_sigmoid.add(Dense(124, activation='relu'))# second layer 124 neurons
model_sigmoid.add(Dense(50, activation='relu'))# third layer 50 neurons
model_sigmoid.add(Dense(output_dim, activation='softmax'))
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

```
model_sigmoid.summary()
```

↗ Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_1 (Dense)	(None, 500)	392500
dense_2 (Dense)	(None, 124)	62124
dense_3 (Dense)	(None, 50)	6250
dense_4 (Dense)	(None, 10)	510
=====	=====	=====
Total params: 461,384		
Trainable params: 461,384		
Non-trainable params: 0		
=====		

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb
```

↗

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:79

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python
Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

60000/60000 [=====] - 7s 113us/step - loss: 1.2421 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 7s 114us/step - loss: 0.4223 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 8s 126us/step - loss: 0.3238 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 6s 104us/step - loss: 0.2812 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2538 - acc: 0.9

Epoch 6/20

60000/60000 [=====] - 6s 95us/step - loss: 0.2321 - acc: 0.9

Epoch 7/20

60000/60000 [=====] - 6s 93us/step - loss: 0.2154 - acc: 0.9

Epoch 8/20

60000/60000 [=====] - 6s 94us/step - loss: 0.2004 - acc: 0.9

Epoch 9/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1874 - acc: 0.9

Epoch 10/20

60000/60000 [=====] - 5s 89us/step - loss: 0.1765 - acc: 0.9

Epoch 11/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1659 - acc: 0.9

Epoch 12/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1568 - acc: 0.9

Epoch 13/20

60000/60000 [=====] - 6s 94us/step - loss: 0.1483 - acc: 0.9

Epoch 14/20

60000/60000 [=====] - 6s 93us/step - loss: 0.1410 - acc: 0.9

Epoch 15/20

60000/60000 [=====] - 5s 89us/step - loss: 0.1338 - acc: 0.9

Epoch 16/20

60000/60000 [=====] - 6s 92us/step - loss: 0.1272 - acc: 0.9

Epoch 17/20

60000/60000 [=====] - 6s 96us/step - loss: 0.1212 - acc: 0.9

Epoch 18/20

60000/60000 [=====] - 6s 98us/step - loss: 0.1158 - acc: 0.9

Epoch 19/20

```
60000/60000 [=====] - 6s 97us/step - loss: 0.1107 - acc: 0.9
Epoch 20/20
60000/60000 [=====] - 6s 94us/step - loss: 0.1059 - acc: 0.9
```

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

➡ Test score: 0.11591833184994757
Test accuracy: 0.9646



▼ 5 layered mlp with relu activation

```
model_relu5 = Sequential()
model_relu5.add(Dense(500, activation='relu', input_shape=(input_dim,)))# first layer 500
model_relu5.add(Dense(324, activation='relu'))# second layer 324 neurons
```

```

model_relu5.add(Dense(250, activation='relu'))# third layer 250 neurons
model_relu5.add(Dense(150, activation='relu'))# forth layer 150 neurons
model_relu5.add(Dense(50, activation='relu'))# fifth layer 50 neurons
model_relu5.add(Dense(output_dim, activation='softmax'))

```

```
model_relu5.summary()
```

↳ Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
dense_5 (Dense)	(None, 500)	392500
dense_6 (Dense)	(None, 324)	162324
dense_7 (Dense)	(None, 250)	81250
dense_8 (Dense)	(None, 150)	37650
dense_9 (Dense)	(None, 50)	7550
dense_10 (Dense)	(None, 10)	510
=====		
Total params: 681,784		
Trainable params: 681,784		
Non-trainable params: 0		

```

model_relu5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
history = model_relu5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbos

```

↳

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 186us/step - loss: 0.2587 - acc: 0
Epoch 2/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0963 - acc: 0
Epoch 3/20
60000/60000 [=====] - 10s 171us/step - loss: 0.0669 - acc: 0
Epoch 4/20
60000/60000 [=====] - 11s 177us/step - loss: 0.0513 - acc: 0
Epoch 5/20
60000/60000 [=====] - 11s 177us/step - loss: 0.0387 - acc: 0
Epoch 6/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0332 - acc: 0
Epoch 7/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0283 - acc: 0
Epoch 8/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0249 - acc: 0
Epoch 9/20
60000/60000 [=====] - 11s 181us/step - loss: 0.0244 - acc: 0
Epoch 10/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0193 - acc: 0
Epoch 11/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0199 - acc: 0
Epoch 12/20
60000/60000 [=====] - 10s 171us/step - loss: 0.0169 - acc: 0
Epoch 13/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0176 - acc: 0
Epoch 14/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0181 - acc: 0
Epoch 15/20
60000/60000 [=====] - 11s 177us/step - loss: 0.0130 - acc: 0
Epoch 16/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0103 - acc: 0
Epoch 17/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0138 - acc: 0
Epoch 18/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0108 - acc: 0
Epoch 19/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0131 - acc: 0
Epoch 20/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0094 - acc: 0
```

```
score = model_relu5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
```




```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

☞ Test score: 0.09293281461207743
Test accuracy: 0.9813



➤ 3 layered mlp with he_initializer without any batch_normalizer and c

```
#he_normal(seed=None)
model_relu = Sequential()
model_relu.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
model_relu.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

☞

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
dense_11 (Dense)	(None, 500)	392500
dense_12 (Dense)	(None, 124)	62124
dense_13 (Dense)	(None, 50)	6250
dense_14 (Dense)	(None, 10)	510
=====		
Total params: 461,384		
Trainable params: 461,384		
Non-trainable params: 0		

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 129us/step - loss: 0.2391 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0889 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0563 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0392 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0306 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0244 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0204 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0179 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0151 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0126 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 7s 114us/step - loss: 0.0153 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0129 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0127 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 7s 115us/step - loss: 0.0085 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0100 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0121 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0071 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0093 - acc: 0.
Epoch 19/20
60000/60000 [=====] - 7s 116us/step - loss: 0.0096 - acc: 0.
Epoch 20/20
60000/60000 [=====] - 7s 118us/step - loss: 0.0090 - acc: 0.
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))
```


```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

➞ Test score: 0.10216689538143828
Test accuracy: 0.9803



➤ 5 layered mlp with initializer = he_normal and relu activation without

```
#he_normal(seed=None)
model_relu5 = Sequential()
model_relu5.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=he_normal(seed=None)))
model_relu5.add(Dense(324, activation='relu', kernel_initializer=he_normal(seed=None)))
model_relu5.add(Dense(250, activation='relu', kernel_initializer=he_normal(seed=None)))
model_relu5.add(Dense(150, activation='relu', kernel_initializer=he_normal(seed=None)))
model_relu5.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)))
model_relu5.add(Dense(output_dim, activation='softmax'))
```

```
model_relu5.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

➞

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 185us/step - loss: 0.2317 - acc: 0
Epoch 2/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0889 - acc: 0
Epoch 3/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0611 - acc: 0
Epoch 4/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0483 - acc: 0
Epoch 5/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0368 - acc: 0
Epoch 6/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0317 - acc: 0
Epoch 7/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0272 - acc: 0
Epoch 8/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0252 - acc: 0
Epoch 9/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0258 - acc: 0
Epoch 10/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0217 - acc: 0
Epoch 11/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0177 - acc: 0
Epoch 12/20
60000/60000 [=====] - 11s 179us/step - loss: 0.0158 - acc: 0
Epoch 13/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0153 - acc: 0
Epoch 14/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0149 - acc: 0
Epoch 15/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0134 - acc: 0
Epoch 16/20
60000/60000 [=====] - 11s 176us/step - loss: 0.0117 - acc: 0
Epoch 17/20
60000/60000 [=====] - 11s 182us/step - loss: 0.0151 - acc: 0
Epoch 18/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0135 - acc: 0
Epoch 19/20
60000/60000 [=====] - 11s 178us/step - loss: 0.0115 - acc: 0
Epoch 20/20
60000/60000 [=====] - 11s 183us/step - loss: 0.0110 - acc: 0
```

```
score = model_relu5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```


```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

➞ Test score: 0.08569526613150633
Test accuracy: 0.9824



➤ 3 layered mlp with batch normalizer without dropout activation is re

```
from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
model_batch.add(BatchNormalization())

model_batch.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
```

➞

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy']
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbos
```

☞ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 160us/step - loss: 0.2196 - acc: 0
Epoch 2/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0815 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0521 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0385 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0326 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 8s 140us/step - loss: 0.0254 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0229 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0189 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0159 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 9s 143us/step - loss: 0.0179 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 9s 145us/step - loss: 0.0179 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0136 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0136 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0123 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 8s 138us/step - loss: 0.0115 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0098 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 9s 142us/step - loss: 0.0110 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0103 - acc: 0.
Epoch 19/20
60000/60000 [=====] - 8s 139us/step - loss: 0.0080 - acc: 0.
Epoch 20/20
60000/60000 [=====] - 8s 136us/step - loss: 0.0060 - acc: 0.
```

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```


```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

➞ Test score: 0.08583917959941027
Test accuracy: 0.9802



➤ 5 layered mlp with batch normalizer without dropout activation is re

```
model_batch = Sequential()

model_batch.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
model_batch.add(BatchNormalization())

model_batch.add(Dense(324, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(250, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(150, activation='relu', input_shape=(input_dim,), kernel_initializer
model_batch.add(BatchNormalization())

model_batch.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_batch.add(BatchNormalization())
```



```
model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=0)
```

☞ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 15s 242us/step - loss: 0.2159 - acc: 0
Epoch 2/20
60000/60000 [=====] - 13s 214us/step - loss: 0.0855 - acc: 0
Epoch 3/20
60000/60000 [=====] - 13s 219us/step - loss: 0.0591 - acc: 0
Epoch 4/20
60000/60000 [=====] - 13s 224us/step - loss: 0.0469 - acc: 0
Epoch 5/20
60000/60000 [=====] - 13s 225us/step - loss: 0.0399 - acc: 0
Epoch 6/20
60000/60000 [=====] - 13s 219us/step - loss: 0.0332 - acc: 0
Epoch 7/20
60000/60000 [=====] - 13s 219us/step - loss: 0.0305 - acc: 0
Epoch 8/20
60000/60000 [=====] - 13s 218us/step - loss: 0.0271 - acc: 0
Epoch 9/20
60000/60000 [=====] - 13s 216us/step - loss: 0.0239 - acc: 0
Epoch 10/20
60000/60000 [=====] - 13s 216us/step - loss: 0.0237 - acc: 0
Epoch 11/20
60000/60000 [=====] - 13s 219us/step - loss: 0.0217 - acc: 0
Epoch 12/20
60000/60000 [=====] - 14s 225us/step - loss: 0.0166 - acc: 0
Epoch 13/20
60000/60000 [=====] - 13s 220us/step - loss: 0.0207 - acc: 0
Epoch 14/20
60000/60000 [=====] - 13s 221us/step - loss: 0.0138 - acc: 0
Epoch 15/20
60000/60000 [=====] - 13s 223us/step - loss: 0.0148 - acc: 0
Epoch 16/20
60000/60000 [=====] - 13s 216us/step - loss: 0.0133 - acc: 0
Epoch 17/20
60000/60000 [=====] - 14s 230us/step - loss: 0.0179 - acc: 0
Epoch 18/20
60000/60000 [=====] - 13s 216us/step - loss: 0.0132 - acc: 0
Epoch 19/20
60000/60000 [=====] - 13s 211us/step - loss: 0.0114 - acc: 0
Epoch 20/20
60000/60000 [=====] - 13s 214us/step - loss: 0.0116 - acc: 0
```

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```


```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

➞ Test score: 0.07765144463920151
Test accuracy: 0.9805



➤ 3 layered mlp without batch normalizer with dropout activation is re

```
from keras.layers import Dropout
```

```
from keras.layers import Dropout
```

```
model_drop = Sequential()
```

```
model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))
```

```
model_drop.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))
```

```
model_drop.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(Dropout(rate=0.5))
```

```
model_drop.add(Dense(output_dim, activation='softmax'))
```

⚠ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1445: tf.nn.softmax_cross_entropy_with_logits is deprecated and will be removed in a future version. Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

⚡ Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 10s 161us/step - loss: 0.8373 - acc: 0.0000
Epoch 2/20
60000/60000 [=====] - 8s 133us/step - loss: 0.3337 - acc: 0.0000
Epoch 3/20
60000/60000 [=====] - 8s 133us/step - loss: 0.2548 - acc: 0.0000
Epoch 4/20
60000/60000 [=====] - 8s 132us/step - loss: 0.2166 - acc: 0.0000
Epoch 5/20
60000/60000 [=====] - 8s 133us/step - loss: 0.1895 - acc: 0.0000
Epoch 6/20
60000/60000 [=====] - 8s 140us/step - loss: 0.1690 - acc: 0.0000
Epoch 7/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1536 - acc: 0.0000
Epoch 8/20
60000/60000 [=====] - 8s 133us/step - loss: 0.1450 - acc: 0.0000
Epoch 9/20
60000/60000 [=====] - 8s 132us/step - loss: 0.1350 - acc: 0.0000
Epoch 10/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1281 - acc: 0.0000
Epoch 11/20
60000/60000 [=====] - 8s 136us/step - loss: 0.1206 - acc: 0.0000
Epoch 12/20
60000/60000 [=====] - 8s 135us/step - loss: 0.1177 - acc: 0.0000
Epoch 13/20
60000/60000 [=====] - 8s 140us/step - loss: 0.1084 - acc: 0.0000
Epoch 14/20
60000/60000 [=====] - 8s 128us/step - loss: 0.1024 - acc: 0.0000
Epoch 15/20
60000/60000 [=====] - 8s 135us/step - loss: 0.0994 - acc: 0.0000
Epoch 16/20
60000/60000 [=====] - 8s 133us/step - loss: 0.0933 - acc: 0.0000
Epoch 17/20
60000/60000 [=====] - 8s 137us/step - loss: 0.0936 - acc: 0.0000
Epoch 18/20
60000/60000 [=====] - 8s 129us/step - loss: 0.0885 - acc: 0.0000
Epoch 19/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0843 - acc: 0.0000
Epoch 20/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0894 - acc: 0.0000
```

```

score(*model_drop.evaluate(X_test, Y_test, verbose=0))
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

➞ Test score: 0.07902954334445494
Test accuracy: 0.9818



➤ 5 layered mlp without batch normalizer with dropout activation is re

```

model_drop = Sequential()

model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,)), kernel_initializer=
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(324, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

```

```

model_drop.add(Dense(250, activation='relu',kernel_initializer=he_normal(seed=None)))
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(150, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
#model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

```

⏏ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose

```

⏏

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:79

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python
Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorfl

60000/60000 [=====] - 14s 232us/step - loss: 1.4113 - acc: 0

Epoch 2/20

60000/60000 [=====] - 13s 212us/step - loss: 0.4782 - acc: 0

Epoch 3/20

60000/60000 [=====] - 13s 212us/step - loss: 0.3456 - acc: 0

Epoch 4/20

60000/60000 [=====] - 13s 211us/step - loss: 0.2901 - acc: 0

Epoch 5/20

60000/60000 [=====] - 13s 213us/step - loss: 0.2567 - acc: 0

Epoch 6/20

60000/60000 [=====] - 13s 210us/step - loss: 0.2308 - acc: 0

Epoch 7/20

60000/60000 [=====] - 13s 214us/step - loss: 0.2145 - acc: 0

Epoch 8/20

60000/60000 [=====] - 13s 213us/step - loss: 0.1999 - acc: 0

Epoch 9/20

60000/60000 [=====] - 13s 214us/step - loss: 0.1896 - acc: 0

Epoch 10/20

60000/60000 [=====] - 13s 213us/step - loss: 0.1829 - acc: 0

Epoch 11/20

60000/60000 [=====] - 13s 209us/step - loss: 0.1721 - acc: 0

Epoch 12/20

60000/60000 [=====] - 13s 214us/step - loss: 0.1607 - acc: 0

Epoch 13/20

60000/60000 [=====] - 13s 212us/step - loss: 0.1567 - acc: 0

Epoch 14/20

60000/60000 [=====] - 13s 214us/step - loss: 0.1556 - acc: 0

Epoch 15/20

60000/60000 [=====] - 13s 213us/step - loss: 0.1442 - acc: 0

Epoch 16/20

60000/60000 [=====] - 13s 212us/step - loss: 0.1396 - acc: 0

Epoch 17/20

60000/60000 [=====] - 13s 214us/step - loss: 0.1433 - acc: 0

Epoch 18/20

60000/60000 [=====] - 13s 209us/step - loss: 0.1344 - acc: 0

Epoch 19/20

```
60000/60000 [=====] - 13s 211us/step - loss: 0.1261 - acc: 0
Epoch 20/20
60000/60000 [=====] - 13s 209us/step - loss: 0.1229 - acc: 0
```

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

```
↳ Test score: 0.09795631028999714
Test accuracy: 0.9796
```



▼ 3 layered mlp with batch normalizer with dropout activation is relu a

```
model_drop = Sequential()

model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,), kernel_initializer=
```

```

model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(124, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose

```

☞ Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 12s 199us/step - loss: 0.7329 - acc: 0
Epoch 2/20
60000/60000 [=====] - 9s 157us/step - loss: 0.3083 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 10s 159us/step - loss: 0.2327 - acc: 0
Epoch 4/20
60000/60000 [=====] - 10s 166us/step - loss: 0.2002 - acc: 0
Epoch 5/20
60000/60000 [=====] - 10s 160us/step - loss: 0.1712 - acc: 0
Epoch 6/20
60000/60000 [=====] - 10s 162us/step - loss: 0.1545 - acc: 0
Epoch 7/20
60000/60000 [=====] - 9s 155us/step - loss: 0.1431 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 10s 161us/step - loss: 0.1303 - acc: 0
Epoch 9/20
60000/60000 [=====] - 10s 169us/step - loss: 0.1284 - acc: 0
Epoch 10/20
60000/60000 [=====] - 10s 166us/step - loss: 0.1183 - acc: 0
Epoch 11/20
60000/60000 [=====] - 10s 161us/step - loss: 0.1094 - acc: 0
Epoch 12/20
60000/60000 [=====] - 10s 159us/step - loss: 0.1069 - acc: 0
Epoch 13/20
60000/60000 [=====] - 9s 157us/step - loss: 0.1029 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 10s 160us/step - loss: 0.0971 - acc: 0
Epoch 15/20
60000/60000 [=====] - 10s 161us/step - loss: 0.0922 - acc: 0
Epoch 16/20
60000/60000 [=====] - 9s 158us/step - loss: 0.0898 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 10s 167us/step - loss: 0.0839 - acc: 0
Epoch 18/20
60000/60000 [=====] - 10s 159us/step - loss: 0.0847 - acc: 0
Epoch 19/20
60000/60000 [=====] - 10s 165us/step - loss: 0.0809 - acc: 0
Epoch 20/20
60000/60000 [=====] - 10s 158us/step - loss: 0.0762 - acc: 0

```



```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoch

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

➞ Test score: 0.06307774550883914
Test accuracy: 0.9816



➤ 5 layered mlp with batch normalizer with dropout activation is relu :

```

model_drop = Sequential()

model_drop.add(Dense(500, activation='relu', input_shape=(input_dim,)), kernel_initializer=
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(324, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())

```

```
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(250, activation='relu', kernel_initializer=he_normal(seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(150, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(50, activation='relu', kernel_initializer=he_normal(seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(rate=0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose
```



Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 14s 225us/step - loss: 0.1271 - acc: 0

Epoch 2/20

60000/60000 [=====] - 13s 215us/step - loss: 0.1220 - acc: 0

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# print(history.history.keys())
```

```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
```

```
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
```

```
# we will get val_loss and val_acc only when you pass the paramter validation_data
```

```
# val_loss : validation loss
```

```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
# for each key in history.history we will have a list of length equal to number of epoch
```

```
vy = history.history['val_loss']
```


```
ty = history.history['loss']
```

```
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.10175361420742388

Test accuracy: 0.983



```
# Please compare all your models using Prettytable library
```

```
from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["layer", "type", "log_loss", "accuracy"]
```

```

x.add_row(['2_layer','without dropout and batch_norm',0.07829980347480105,0.9832])
x.add_row(['2_layer','without dropout and with batch_norm',0.0779731295770489,0.982])
x.add_row(['2_layer','with dropout and without batch_norm',0.06712127507654286,0.9829])
x.add_row(['2_layer','with dropout and with batch_norm',0.057800181532799616,0.9826])
x.add_row(['3_layer','without dropout and batch_norm',0.10216689538143828,0.9803])
x.add_row(['5_layer','without dropout and batch_norm',0.08569526613150633,0.9824])
x.add_row(['3_layer','without dropout and with batch_norm',0.08583917959941027,0.9802])
x.add_row(['5_layer','without dropout and with batch_norm',0.07765144463920151,0.9805])
x.add_row(['3_layer','with dropout and without batch_norm',0.07902954334445494,0.9818])
x.add_row(['5_layer','with dropout and without batch_norm',0.09795631028999714,0.9796])
x.add_row(['3_layer','with dropout and with batch_norm',0.06307774550883914,0.9816])
x.add_row(['5_layer','with dropout and with batch_norm',0.10175361420742388,0.983])

```

```
print(x)
```

	layer	type	log_loss	accuracy
	2_layer	without dropout and batch_norm	0.07829980347480105	0.9832
	2_layer	without dropout and with batch_norm	0.0779731295770489	0.982
	2_layer	with dropout and without batch_norm	0.06712127507654286	0.9829
	2_layer	with dropout and with batch_norm	0.057800181532799616	0.9826
	3_layer	without dropout and batch_norm	0.10216689538143828	0.9803
	5_layer	without dropout and batch_norm	0.08569526613150633	0.9824
	3_layer	without dropout and with batch_norm	0.08583917959941027	0.9802
	5_layer	without dropout and with batch_norm	0.07765144463920151	0.9805
	3_layer	with dropout and without batch_norm	0.07902954334445494	0.9818
	5_layer	with dropout and without batch_norm	0.09795631028999714	0.9796
	3_layer	with dropout and with batch_norm	0.06307774550883914	0.9816
	5_layer	with dropout and with batch_norm	0.10175361420742388	0.983